

## Problem 12: PodSecurity Policy

- Enable Pod Security policy on the cluster
- Create a PodSecurityPolicy demo to prevent the creation of all privileged pods.
- Create a Role and Role binding to use this policy and assign it to a new service account sam. Do this in a new test namespace.

### Solution 12:

Pod Security Policies enable fine-grained authorization of pod creation and updates. PodSecurityPolicy is deprecated in v1.21+ and will be unavailable in v1.25+ but it is still part of CKS certification, so keeping this scenario as it is.

I will be using my killercoda Kubernetes playground(<https://killercoda.com/saiyampathak/scenario/kube124>) for this example and will be using the same example as mentioned in the kubernetes documentation

First enable the PodSecurityPolicy plugin by changing the `/etc/kubernetes/manifests/kube-apiserver.yaml` file .

```
spec:
  containers:
  - command:
    - kube-apiserver
    - --advertise-address=172.17.0.39
    - --allow-privileged=true
    - --authorization-mode=Node,RBAC
    - --client-ca-file=/etc/kubernetes/pki/ca.crt
    - --enable-admission-plugins=NodeRestriction,PodSecurityPolicy
```

Create a PodSecurityPolicy to deny privileged pods

Note - Before enabling the admission plugin make sure to have a default allow policy so that pods are not prevented

```
apiVersion: policy/v1beta1
kind: PodSecurityPolicy
```

```

metadata:
  name: demo
spec:
  privileged: false # Don't allow privileged pods!
  # The rest fills in some required fields.
  selinux:
    rule: RunAsAny
  supplementalGroups:
    rule: RunAsAny
  runAsUser:
    rule: RunAsAny
  fsGroup:
    rule: RunAsAny
  volumes:
  - '*'

```

If you now try to create the privileged pod, it will fail

```

apiVersion: v1
kind: Pod
metadata:
  name: test-pod-1
  namespace: default
spec:
  containers:
  - name: centos
    image: centos
    command: ['sh', '-c', 'sleep 999']
    securityContext:
      privileged: true

```

```

Error from server (Forbidden): error when creating "po": pods
"test-pod-1" is forbidden: PodSecurityPolicy: unable to admit
pod: [spec.containers[0].securityContext.privileged: Invalid
value: true: Privileged containers are not allowed]

```

Create service account in test namespace

```
kubectl create ns test
namespace/test created

kubectl create sa sam -n test
serviceaccount/sam created
```

Create a Role and RoleBinding and bind it to SA sam

```
kubectl create role psp --verb=use
--resource=podsecuritypolicy --resource-name=demo -n test
role.rbac.authorization.k8s.io/psp created

kubectl create -n test rolebinding psp-binding --role=psp
--serviceaccount=test:sam
rolebinding.rbac.authorization.k8s.io/psp-binding created
```

No there can be scenarios where you can be asked to create a specific policy and you can control all of the below

Running of privileged containers	<a href="#">privileged</a>
Usage of host namespaces	<a href="#">hostPID</a> , <a href="#">hostIPC</a>
Usage of host networking and ports	<a href="#">hostNetwork</a> , <a href="#">hostPorts</a>
Usage of volume types	<a href="#">volumes</a>
Usage of the host filesystem	<a href="#">allowedHostPaths</a>
Allow specific FlexVolume drivers	<a href="#">allowedFlexVolumes</a>
Allocating an FSGroup that owns the pod's volumes	<a href="#">fsGroup</a>
Requiring the use of a read only root file system	<a href="#">readOnlyRootFilesystem</a>
The user and group IDs of the container	<a href="#">runAsUser</a> , <a href="#">runAsGroup</a> , <a href="#">supplementalGroups</a>
Restricting escalation to root privileges	<a href="#">allowPrivilegeEscalation</a> , <a href="#">defaultAllowPrivilegeEscalation</a>
Linux capabilities	<a href="#">defaultAddCapabilities</a> , <a href="#">requiredDropCapabilities</a> , <a href="#">allowedCapabilities</a>
The SELinux context of the container	<a href="#">seLinux</a>
The Allowed Proc Mount types for the container	<a href="#">allowedProcMountTypes</a>
The AppArmor profile used by containers	<a href="#">annotations</a>
The seccomp profile used by containers	<a href="#">annotations</a>
The sysctl profile used by containers	<a href="#">forbiddenSysctls</a> , <a href="#">allowedUnsafeSysctls</a>

Bookmark this page -

<https://kubernetes.io/docs/concepts/policy/pod-security-policy/>

You can find the examples here.

**Example**, if you are asked to create a PodSecurityPolicy that allows hostPath volumes only for directory /foo then you can use the below snippet.

```
You can add the following to the PodSecurityPolicy
allowedHostPaths:
  # This allows "/foo", "/foo/", "/foo/bar" etc., but
  # disallows "/foo1", "/etc/foo" etc.
  # "/foo/.." is never valid.
  - pathPrefix: "/foo"
    readOnly: true # only allow read-only mounts
```

This snippet is directly taken from the Kubernetes documentation so make sure you understand the concepts and bookmark the page.